

Family	Description	Common Series	Nigerian Real-World Use Case
General Purpose	Balanced compute, memory, and networking for everyday applications. Good for websites, APIs, small databases, and business apps.	M, T	A Nigerian fintech startup hosting its customer dashboard and REST APIs for mobile banking.
Compute Optimized	High CPU performance for compute-heavy workloads like analytics, gaming servers, transcoding, and HPC.	C	A Lagos media company processing large volumes of video encoding for online streaming.
Memory Optimized	Large RAM capacity for workloads needing fast in-memory processing such as Redis, SAP, analytics, and caching.	R, X	A Nigerian bank running fraud detection and in-memory transaction analytics using Redis.
Storage Optimized	High-speed local SSD/HDD storage with very high IOPS for databases, big data, Elasticsearch, and streaming.	I, D	A telecom company in Nigeria storing and analyzing massive call-detail records and logs.
Accelerated Computing	Uses GPUs or hardware accelerators for machine learning, AI, graphics rendering, and scientific computing.	P, G	A Nigerian AI startup training computer vision models for traffic monitoring in Abuja and Lagos.

Instance Type	Family	vCPUs	RAM (GiB)	Network	Use Case	Monthly Cost (us-east-1, Linux On-Demand)*	Right Instance For
t3.micro	General Purpose (T)	2	1	Up to 5 Gbps	Small websites and testing	~\$7.59/month	Personal blog or startup MVP
m5.large	General Purpose (M)	2	8	Up to 10 Gbps	Balanced business applications	~\$69.12/month	Production web applications and APIs
c5.xlarge	Compute Optimized (C)	4	8	Up to 10 Gbps	CPU-intensive workloads	~\$124.10/month	Video transcoding or analytics engine
r5.2xlarge	Memory Optimized (R)	8	64	Up to 10 Gbps	In-memory databases and caching	~\$504.43/month	Redis or SAP HANA workloads
i3.large	Storage Optimized (I)	2	15.25	Up to 10 Gbps	High IOPS storage workloads	~\$124.83/month	Elasticsearch or NoSQL databases
g4dn.xlarge	Accelerated Computing (G)	4	16	Up to 25 Gbps	GPU-based AI and graphics	~\$379.61/month	Machine learning inference or rendering

*Monthly estimates are based on approximately 730 hours/month in **us-east-1 Linux On-Demand pricing**.